

Introduction to Intelligent Systems - Lab 5

Tom Apol (s2701650)
Nielis Brouwer (s3850706)
Group: 23

October 15, 2018

Contents

1	Introduction	2
2	Problem Description	2
3	Results	2
4	Discussion	3
5	Conclusion	4
6	Appendix	5
6.1	Figures	5
6.1.1	Prototypes over time	5
6.1.2	Quantisation error over time	11
6.2	Code	17

1 Introduction

In this report we will describe the effects of different amount of prototypes (K) and learning rates (η) with regards to Vector Quantisation of a 2D data set. This algorithm first randomly initialises the prototypes, after which their position is corrected by data points ‘pulling’ the closest prototypes closer to themselves.

The sets used for the experiments are 2D data sets with 1000x2 values where the first column indicates an x-value and the second column indicates a y-value. For our measurements, we have chosen data set `x` (or `w6_1x.mat` to be precise).

2 Problem Description

The problem as mentioned in the assignment is: “Implement Winner-Takes-All unsupervised competitive learning Vector Quantisation, and apply this method to the provided data sets using the Euclidean distance as a measure.” Perform experiments for different number of prototypes K ($K=2$, $K=4$) and different learning rates η . The learning quality is represented by the quantisation error, the lower the error the “better” the prototypes represent the cluster of data points they are closest to.

3 Results

The following algorithm was used to compute the quantisation error, where $d_j^\mu = |\xi^\mu - w_j|$ denotes the Euclidean distance between data point μ and prototype j :

$$H_{VQ} = \sum_{j=1}^K \underbrace{\sum_{\mu=1}^P |\xi^\mu - w_j|}_{\text{data } d_j^\mu} \underbrace{\prod_{k \neq j}^K \Theta(d_k^\mu - d_j^\mu)}_{w_j \text{ is the winner !}}$$

For the results (= plots), see Appendix 6.1. With regards to the code: the code is split into several parts:

- Loading and initialising of variables and data structures: lines 1 - 41
- Computing the quantisation error (H_{VQ}) and performing the algorithm: lines 43 - 94
- Plotting both the prototypes over time (and data points), as well as the quantisation error over time: lines 96 - 169

What we can observe from our plots is the following: when run with a relatively high η (e.g. $\eta = 0.1$), the prototypes move in a rather volatile manner, where

outliers in particular seem to have a relatively large influence on their position, as seen in figures 1 and 4. When multiple prototypes reside in the same large cluster, this can result in the prototypes violently ‘pushing’ each other away, as seen in figure 4. With lower learning rates, both of these effects seem to lessen (see figures 2 and 5).

This is reflected by the quantisation error over time, where the quantisation error ‘stabilises’ early, but with higher values of η , there is still a lot of fluctuation (see figures 7, 8, 10 and 11).

However, with a learning rate that is very low (e.g. $\eta = 0.0001$), we see that the quantisation error becomes rather stable, yet it takes more time to stabilise. This is to be expected, as the step size (= learning rate) is rather low. (see figures 9 and 12). This is also well illustrated by figure 3, where P_1 only travels in a straight line towards its stabilisation point.

Furthermore, multiple prototypes fighting over the same cluster seem to influence each other less (see figure 6).

4 Discussion

Regarding the question ‘what is a reasonable value for t_{\max} , such that the quantisation error has stabilised at that point?’: this depends on η and K . Furthermore, the stabilisation point is influenced by the initialisation of the prototypes, since a prototype that has been initialised further from its ‘stabilised’ position will take more epochs to get there. For $0.1 \geq \eta \geq 0.01$ and data set x , we can say that a t_{\max} of around 20 seems alright. In figures 7, 8, 10 and 11 the stabilisation point generally lies in front of that point, but this is because the initial prototypes lie relatively close to their ‘stabilised’ positions, as seen in their respective ‘Prototypes over time’ figures 1, 2, 4 and 5. To account for this, we chose a t_{\max} that is higher than the ‘stabilisation point’ in our plots.

For $\eta \approx 0.0001$ we can state that it seems to stabilise around 100 epochs or perhaps even later, dependent on K , as seen in figures 9 and 12. The reason for this lower stabilisation rate is the fact that a learning rate of this low magnitude slows down the learning significantly.

This brings us to the next question of interest: “What happens if η is too large or too small?”. As discussed in the ‘Results’ section, for an η that is too large, the quantisation error ‘stabilises’ relatively early, but isn’t very stable. This becomes especially evident when comparing figures 10 and 11. This is due to the prototypes being significantly influenced by the other data points (especially outliers) in their cluster, even after having approached their stabilisation point.

For an η that is too small, the quantisation error becomes very stable, but stabilises after a rather large amount of epochs. This becomes evident when comparing figures 8 and 9, and figures 11 and 12.

Finally we will answer the question "How does the final value of the cost function change with η ?". For this we will take a look at figures 7 to 12. The final (stabilised) value of the cost function/quantisation error seems to slightly decrease when η decreases. However, since figure 12 doesn't seem to fully stabilise yet at our t_{\max} of 100, it is difficult to tell whether this holds for $K=4$.

5 Conclusion

The following conclusions can be summarised from the data presented above:

- We found that the value of the initial point of the prototype has great influence when determining a reasonable value for t_{\max} . For data set x , we deem $\eta \approx 0.01$ to be a reasonable learning rate, and keeping the previous statement in mind, we present $t_{\max} = 20$ as an appropriate corresponding stabilisation point.
- The quantisation error stabilises late when the learning rate is relatively low, but the learning rate becomes very stable.
- A high learning rate results in a early stabilisation, but the quantisation error (and prototype location for that matter) keeps fluctuating around a 'stabilisation point'.
- The final value of the cost function seems to slightly decrease with a lower η .

6 Appendix

6.1 Figures

6.1.1 Prototypes over time

Colours are used to separate our prototypes and their paths, namely red, green, blue and magenta. The initial point of each prototype has been made lighter than the rest for readability purposes. The progression of prototypes is signified by a change in colour from a relatively dark to relatively bright version of their respective colour.

The variables K and η mentioned below represent the number of prototypes and the learning rate respectively. Furthermore, the data used comes from the provided `w6_1.mat` data set.




Figure 1: Prototypes over time with $K=2$ and $\eta=0.1$




Figure 2: Prototypes over time with $K=2$ and $\eta=0.01$




Figure 3: Prototypes over time with $K=2$ and $\eta=0.0001$




Figure 4: Prototypes over time with $K=4$ and $\eta=0.1$




Figure 5: Prototypes over time with $K=4$ and $\eta=0.01$

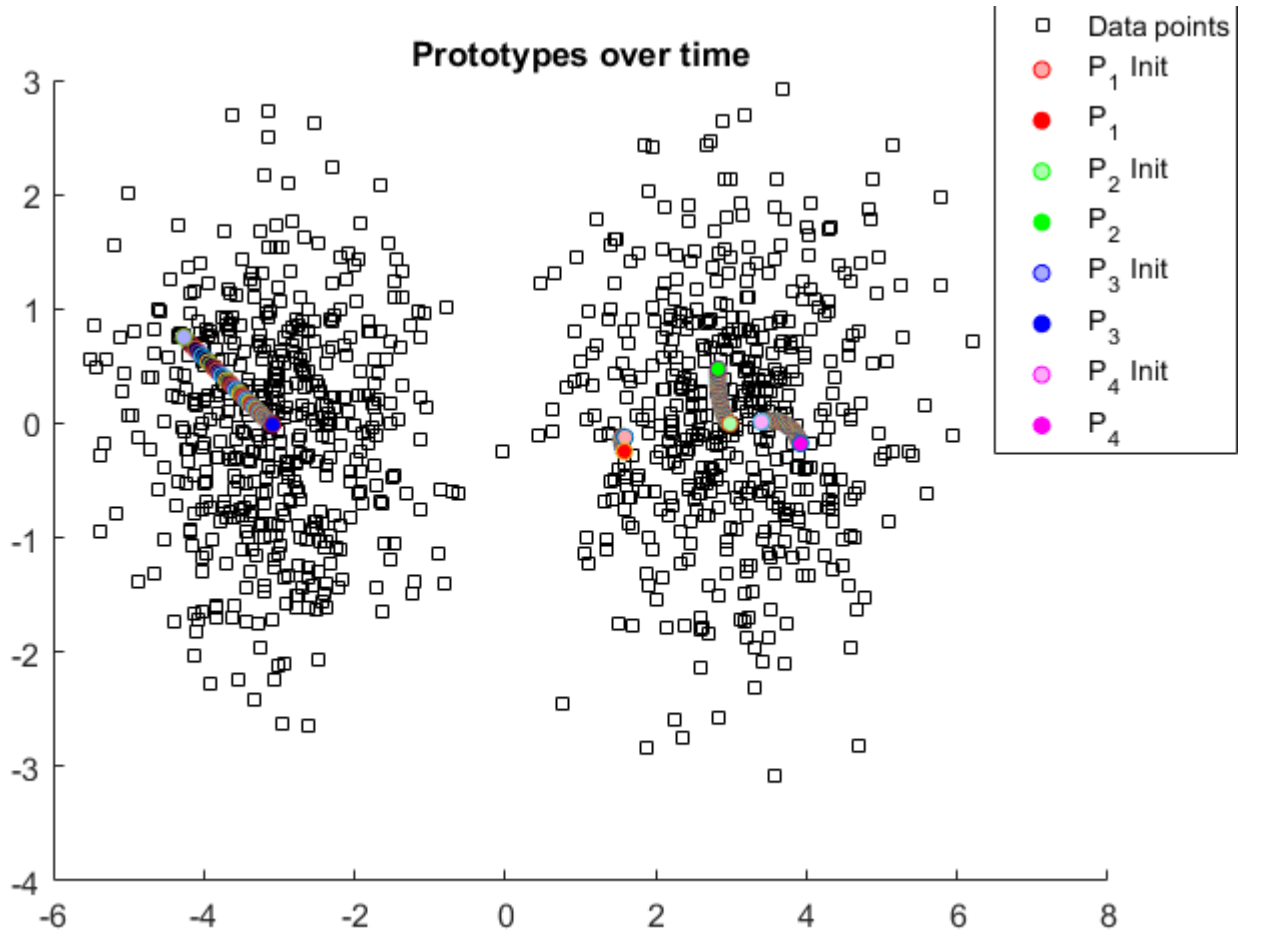


Figure 6: Prototypes over time with $K=4$ and $\eta=0.0001$

6.1.2 Quantisation error over time

The quantisation error has been normalised, such that it has no relation to the size of the dataset. This way we will be able to compare data sets of differing sizes, should we be so inclined.

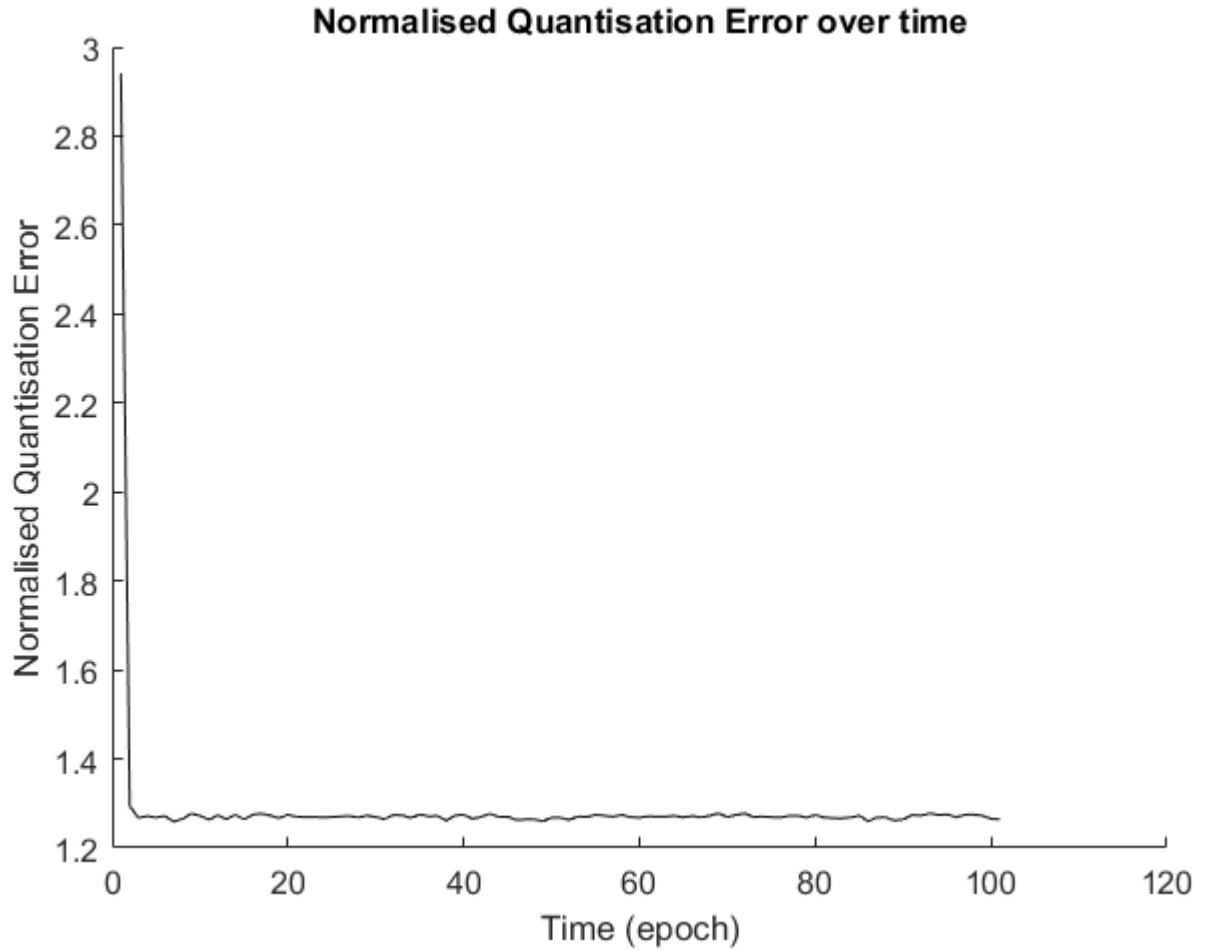


Figure 7: Normalised quantisation error over time with $K=2$ and $\eta=0.1$

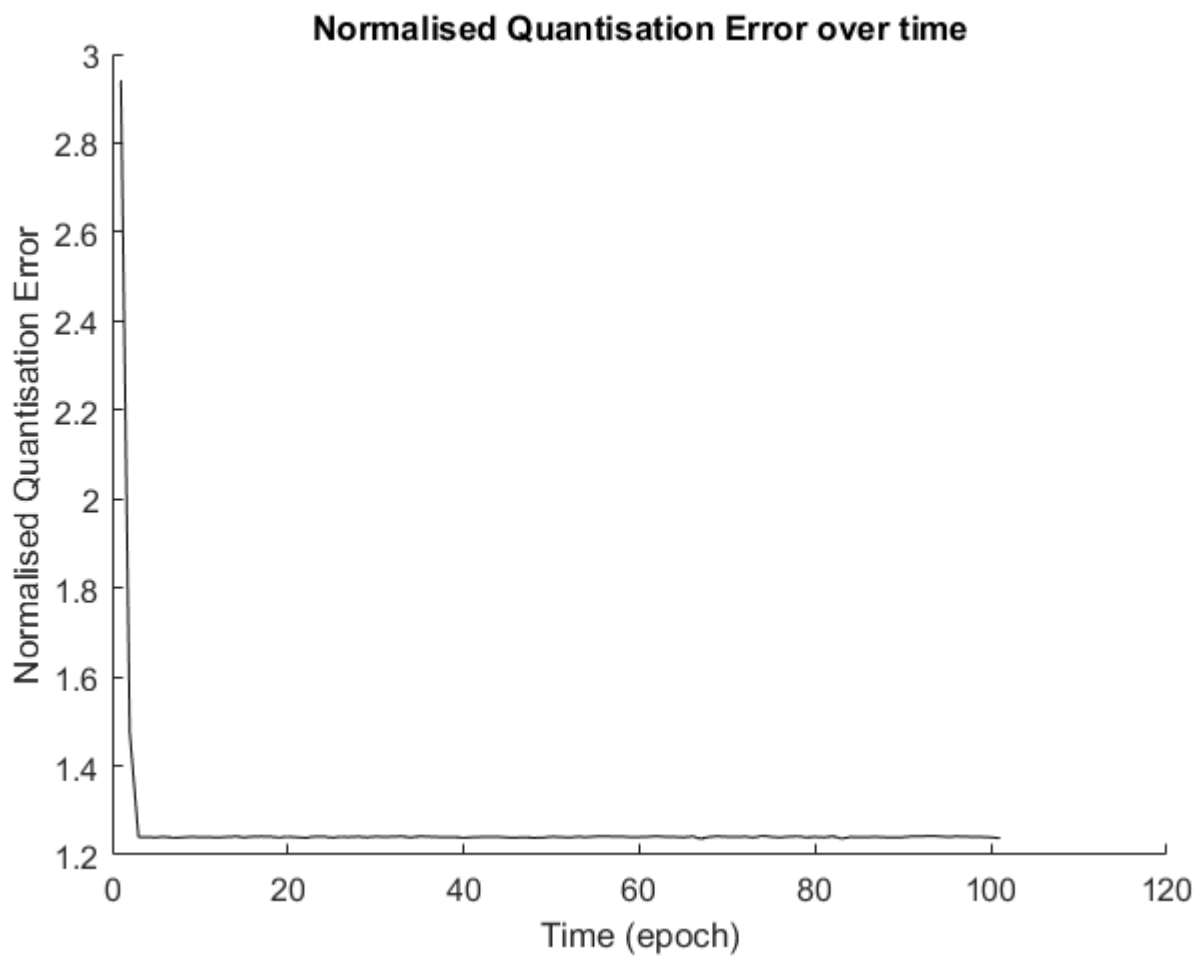


Figure 8: Normalised quantisation error over time with $K=2$ and $\eta=0.01$




Figure 9: Normalised quantisation error over time with $K=2$ and $\eta=0.0001$




Figure 10: Normalised quantisation error over time with $K=4$ and $\eta=0.1$

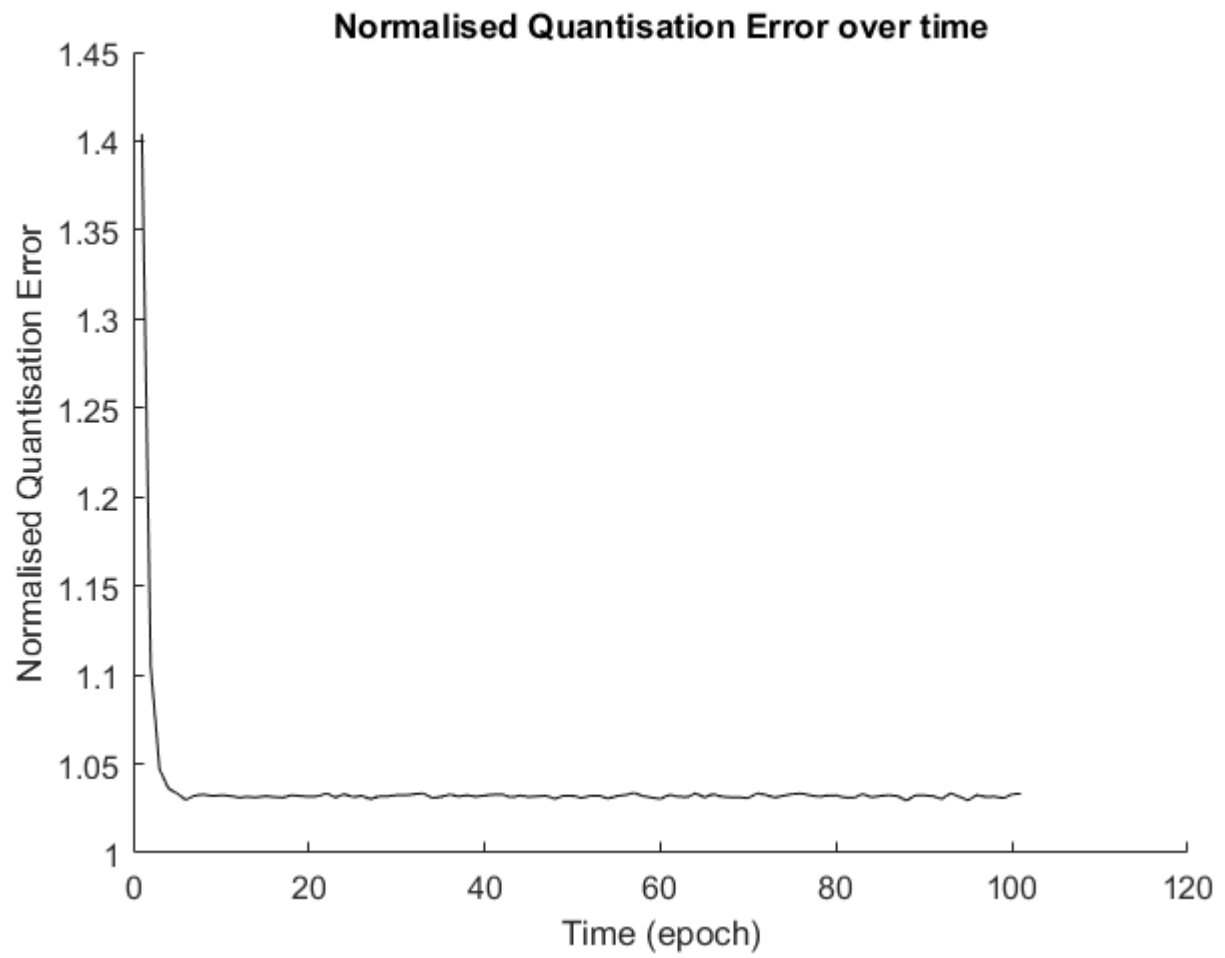


Figure 11: Normalised quantisation error over time with $K=4$ and $\eta=0.01$

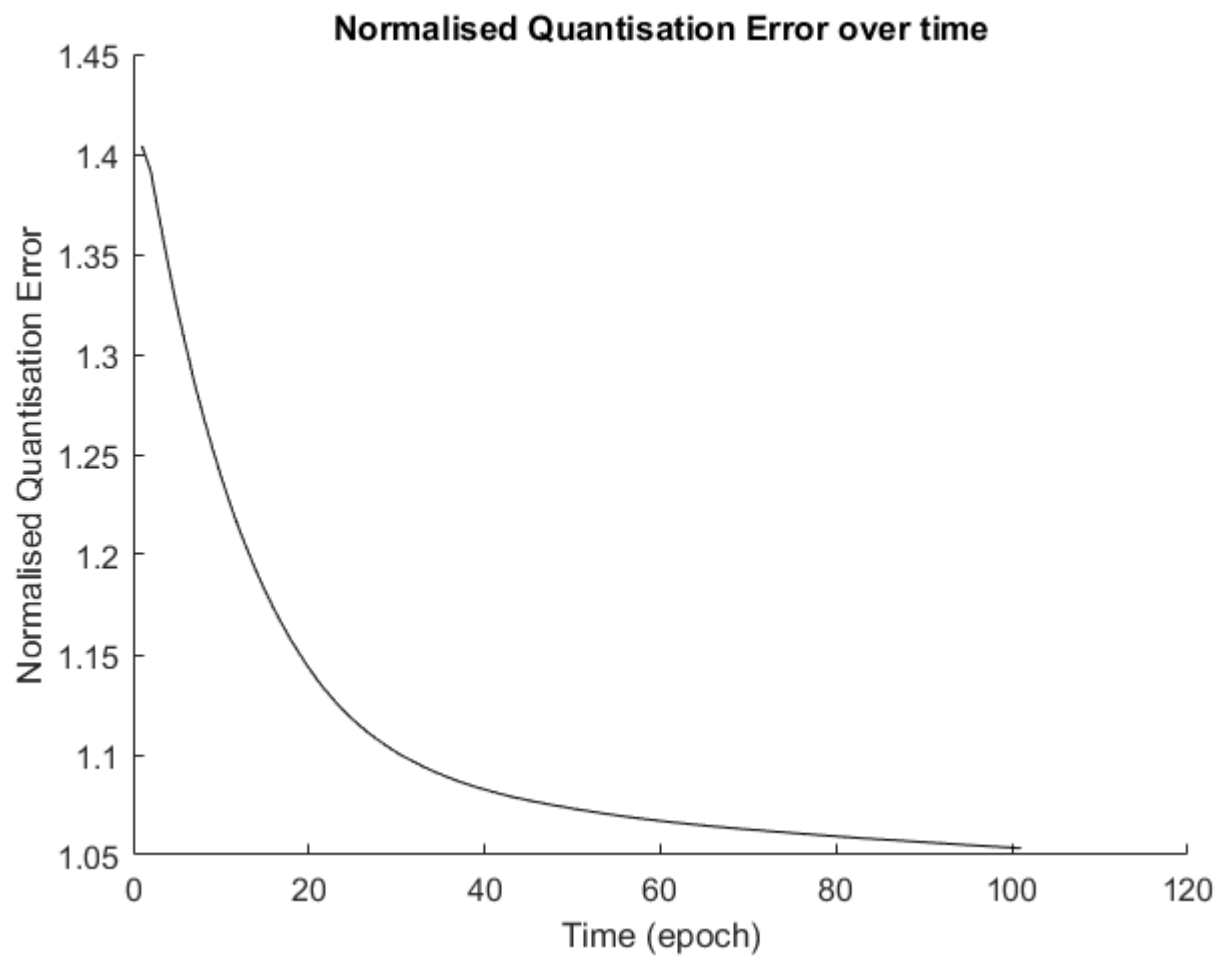


Figure 12: Normalised quantisation error over time with $K=4$ and $\eta=0.0001$

6.2 Code

```
1  clear all;
2  close all;
3
4  % load data
5  data_x = load('w6_1x.mat');
6  data_x = data_x.w6_1x;
7  data_y = load('w6_1y.mat');
8  data_y = data_y.w6_1y;
9  data_z = load('w6_1z.mat');
10 data_z = data_z.w6_1z;
11
12 data = data_x; % replace with dataset of your choice.
13
14 % Initialise RNG
15 rng('default');
16
17 % set colour rgb values
18 red = [1 0 0];
19 green = [0 1 0];
20 blue = [0 0 1];
21 magenta = [1 0 .93];
22 white = [1 1 1];
23 colour_init_1 = (red + white*2)./3;
24 colour_init_2 = (green + white*2)./3;
25 colour_init_3 = (blue + white*2)./3;
26 colour_init_4 = (magenta + white*2)./3;
27
28 % Initialise variables.
29 [P, N] = size(data);
30 K = 4;
31 learning_rate = 0.1;
32 t_max = 100;
33
34 prototypes_over_time = zeros(t_max, K, N);
35 H_VQ_over_time = zeros(t_max+1, 1);
36
37 % Initialise prototypes
38 prototypes = zeros(K, N);
39 for i = 1:K
40     prototypes(i, :) = data(randi(P), :);
41 end
42
43 % Compute initial HVQ
44 for j = randperm(P)
```

```

45     example = data(j, :);
46
47     winner_distance = intmax;
48     for k = 1:K
49         prototype = prototypes(k, :);
50         distance = pdist([example; prototype]);
51         if distance < winner_distance
52             winner_distance = distance;
53         end
54     end
55     H_VQ_over_time(1) = H_VQ_over_time(1) + winner_distance/P;
56 end
57
58 % Permutations
59 for t = 1:t_max
60     H_VQ = 0; % initialise quantisation error for manual sum.
61
62     % Shuffle data set and perform permutations
63     for j = randperm(P)
64         example = data(j, :);
65
66         % initialise winner
67         winner = 1;
68         winner_distance = intmax; % represents infinity
69
70         % get closest prototype
71         for k = 1:K % take rows
72             prototype = prototypes(k, :);
73             distance = pdist([example; prototype]);
74             if distance < winner_distance
75                 winner_distance = distance;
76                 winner = k;
77             end
78         end
79
80         % update winner
81         prototypes(winner, :) = prototypes(winner, :) + ...
82             learning_rate * (example - prototypes(winner, :));
83
84         % update (partial) sum of quantisation error
85         H_VQ = H_VQ + winner_distance;
86     end
87
88     % store prototypes and quantisation error for this iteration
89     prototypes_over_time(t, :, :) = prototypes;
90     H_VQ_over_time(t+1) = H_VQ;

```

```

91
92     % normalise H_VQ
93     H_VQ_over_time(t+1) = H_VQ_over_time(t+1)/P;
94 end
95
96 % plot prototypes and data over time
97 % note that this code is for K = 2 and K = 4
98 figure; hold on;
99
100 % plot data points
101 data_plot = plot(data(:, 1), data(:, 2), 'ksq');
102
103 % plot next iterations with an increasingly coloured line
104 for t = 2:t_max-1
105     plot(prototypes_over_time(t, 1, 1), prototypes_over_time(t, 1, 2), ...
106         'o', 'MarkerFaceColor', [t/t_max /2 + .25 0 0]);
107     plot(prototypes_over_time(t, 2, 1), prototypes_over_time(t, 2, 2), ...
108         'o', 'MarkerFaceColor', [0 t/t_max /2 + .25 0]);
109 end
110
111 % plot first iteration in a whiter form
112 plot(prototypes_over_time(1, 1, 1), prototypes_over_time(1, 1, 2), ...
113     'o', 'MarkerFaceColor', colour_init_1);
114 plot(prototypes_over_time(1, 2, 1), prototypes_over_time(1, 2, 2), ...
115     'o', 'MarkerFaceColor', colour_init_2);
116
117 % plot final iteration with fully coloured line
118 plot(prototypes_over_time(t_max, 1, 1), ...
119     prototypes_over_time(t_max, 1, 2), 'o', 'MarkerFaceColor', red);
120 plot(prototypes_over_time(t_max, 2, 1), ...
121     prototypes_over_time(t_max, 2, 2), 'o', 'MarkerFaceColor', green);
122
123 % hacky way to add extra plots for when K=4
124 if K == 4
125     for t = 2:t_max-1
126         plot(prototypes_over_time(t, 3, 1), prototypes_over_time(t, 3, 2), ...
127             'o', 'MarkerFaceColor', [0 0 t/t_max /2 + .25]);
128         plot(prototypes_over_time(t, 4, 1), prototypes_over_time(t, 4, 2), ...
129             'o', 'MarkerFaceColor', [t/t_max/2+.38 0 t/t_max/2+.33]);
130     end
131
132     plot(prototypes_over_time(1, 3, 1), prototypes_over_time(1, 3, 2), ...
133         'o', 'MarkerFaceColor', colour_init_3);
134     plot(prototypes_over_time(1, 4, 1), prototypes_over_time(1, 4, 2), ...
135         'o', 'MarkerFaceColor', colour_init_4);
136

```

```

137     plot(prototypes_over_time(t_max, 3, 1), ...
138           prototypes_over_time(t_max, 3, 2), 'o', 'MarkerFaceColor', blue);
139     plot(prototypes_over_time(t_max, 4, 1), ...
140           prototypes_over_time(t_max, 4, 2), 'o', 'MarkerFaceColor', magenta);
141 end
142
143 % title and legend
144 title('Prototypes_over_time');
145 legend_plots = zeros(2*K+1, 1);
146 legend_plots(1) = data_plot;
147 legend_plots(2) = plot(nan, nan, 'ro', 'MarkerFaceColor', colour_init_1);
148 legend_plots(3) = plot(nan, nan, 'ro', 'MarkerFaceColor', red);
149 legend_plots(4) = plot(nan, nan, 'go', 'MarkerFaceColor', colour_init_2);
150 legend_plots(5) = plot(nan, nan, 'go', 'MarkerFaceColor', green);
151 if K == 2
152     legend(legend_plots, 'Data_points', 'P_1_Init', 'P_1', ...
153           'P_2_Init', 'P_2');
154 else
155     legend_plots(6) = plot(nan, nan, 'bo', 'MarkerFaceColor', colour_init_3);
156     legend_plots(7) = plot(nan, nan, 'bo', 'MarkerFaceColor', blue);
157     legend_plots(8) = plot(nan, nan, 'mo', 'MarkerFaceColor', colour_init_4);
158     legend_plots(9) = plot(nan, nan, 'mo', 'MarkerFaceColor', magenta);
159     legend(legend_plots, 'Data_points', 'P_1_Init', 'P_1', ...
160           'P_2_Init', 'P_2', 'P_3_Init', 'P_3', ...
161           'P_4_Init', 'P_4');
162 end
163
164 % plot quantisation error over time
165 figure; hold on;
166 plot(1:t_max+1, H_VQ_over_time, 'k-');
167 ylabel('Normalised_Quantisation_Error');
168 xlabel('Time_(epoch)');
169 title('Normalised_Quantisation_Error_over_time');

```